Deep Learning

Lab 5

Value-based Reinforcement Learning

DQN, Atari Pong, and Enhanced DQN

Name : 陳映宬
Department : Electrical Engineering
Student ID : B11107027

May 2026

Contents

1 Introduction **2**

2 Implementation **3**
2.1 Required Implementation Questions . . . 3
2.2 Shared DQN Components . . . 3
2.3 Task 1: CartPole Vanilla DQN . . . 3
2.4 Task 2: Atari Pong Vanilla DQN . . . 4
2.5 Task 3: Enhanced DQN . . . 5
2.5.1 Double DQN . . . 5
2.5.2 Prioritized Experience Replay . . . 6
2.5.3 Multi-step Return . . . 6
2.5.4 Dueling Network . . . 6

3 Experimental Setup and Reproducibility **8**
3.1 Environment . . . 8
3.2 Training Commands . . . 8
3.3 Evaluation Commands . . . 9
3.4 Submitted Checkpoints . . . 10

4 Analysis and Discussion **11**
4.1 Task 1 Result Analysis . . . 11
4.2 Task 2 Result Analysis . . . 12
4.3 Task 3 Main Result . . . 13
4.4 Ablation Study . . . 16

5 Bonus **22**
5.1 Dueling Network . . . 22
5.2 Optimizer and Exploration Tuning . . . 22

6 Conclusion **23**

1 Introduction

This lab studies value-based reinforcement learning using Deep Q-Networks (DQN). DQN combines Q-learning with a neural network function approximator, experience replay, and a target network, allowing the agent to learn action values from high-dimensional observations.

The report covers three tasks. Task 1 implements a vanilla DQN agent for CartPole. Task 2 extends the agent to Atari Pong by using image preprocessing and a convolutional Q-network. Task 3 improves the Pong agent with enhanced DQN techniques, including Double DQN, Prioritized Experience Replay (PER), multi-step returns, and a dueling network architecture.

Key Findings:

- **Task 1:** The final CartPole vanilla DQN snapshot achieves the maximum evaluation score, with average reward **500.00** over 20 consecutive testing seeds.
- **Task 2:** The final vanilla Pong DQN snapshot reaches an average reward of **19.05** over 20 consecutive testing seeds.
- **Task 3:** The final enhanced Pong agent reaches an average reward of **19.10** at 600k environment steps and remains above 19 for all required later milestones.
- The best current configuration is **Dueling Double DQN + PER + 3-step return + exponential epsilon decay + learning-rate decay after 600k steps**.

2 Implementation

2.1 Required Implementation Questions

Table 1 summarizes the required implementation questions specified in the lab description and where they are answered in this report.

Table 1: Required implementation questions and corresponding report sections.

Required item	**Where it is answered**
How to obtain the Bellman error for DQN	Shared DQN Components
How to modify DQN to Double DQN	Task 3: Enhanced DQN – Double DQN
How to implement the memory buffer for PER	Task 3: Enhanced DQN – Prioritized Experience Replay
How to modify the 1-step return to multi-step return	Task 3: Enhanced DQN – Multi-step Return
How Weight & Biases is used to track model performance	Task 2 W&B Tracking and Experimental Setup

2.2 Shared DQN Components

The DQN agent learns an action-value function $Q(s, a; \theta)$ using the Bellman target. For vanilla DQN, the target value is

$$y = r + \gamma(1 - d) \max_{a'} Q_{\text{target}}(s', a'; \theta^-),$$

where d indicates whether the transition is terminal. The network is trained by minimizing the mean squared Bellman error:

$$\mathcal{L}(\theta) = \mathbb{E}\left[(Q(s, a; \theta) - y)^2\right].$$

Experience replay is used to reduce temporal correlation between consecutive samples. Transitions are stored in a replay buffer and mini-batches are sampled for gradient updates. A target network is updated periodically by copying the online Q-network parameters, which stabilizes the bootstrapping target.

Exploration Strategy:

The agent uses an epsilon-greedy policy. With probability ϵ, the agent selects a random action; otherwise, it selects the action with the maximum predicted Q-value:

$$a_t = \begin{cases} \text{random action}, & \text{with probability } \epsilon, \\ \arg\max_a Q(s_t, a; \theta), & \text{otherwise.} \end{cases}$$

2.3 Task 1: CartPole Vanilla DQN

In Task 1, the objective is to keep the pole balanced for as long as possible. The CartPole observation consists of cart position, cart velocity, pole angle, and pole angular velocity. The action space contains two actions: moving the cart left or right.

Q-Network Architecture:

The CartPole agent uses a multi-layer perceptron (MLP) as the Q-network.

Table 2: Task 1 Q-network architecture.

Layer	**Output Size**	**Activation**
Input	4	–
Linear	128	ReLU
Linear	128	ReLU
Linear	2	–

Hyperparameters:

Table 3: Task 1 training hyperparameters.

Parameter	**Value**
Learning rate	1.0×10^{-3}
Batch size	64
Replay buffer size	50k
Discount factor γ	0.99
Epsilon schedule	exponential decay 0.995, $\epsilon_{\min} = 0.01$
Replay start size	1000
Target update frequency	200

2.4 Task 2: Atari Pong Vanilla DQN

Task 2 applies DQN to Atari Pong, where the input is a raw image frame. Since raw frames are high-dimensional and contain redundant visual information, each frame is preprocessed before being passed to the Q-network.

Atari Preprocessing:

The preprocessing pipeline includes:

- converting the frame to grayscale,
- resizing the frame to 84×84,
- stacking the most recent 4 frames as the state representation.

CNN Q-Network Architecture:

The Pong agent uses a convolutional Q-network. The input shape is $4 \times 84 \times 84$.

Table 4: Task 2 CNN Q-network architecture.

Layer	**Channels / Units**	**Kernel / Stride**	**Activation**
Conv2D	32	8×8 / 4	ReLU
Conv2D	64	4×4 / 2	ReLU
Conv2D	64	3×3 / 1	ReLU
Fully Connected	512	–	ReLU
Fully Connected	*[num actions]*	–	–

Hyperparameters:

Table 5: Task 2 training hyperparameters.

Parameter	**Value**
Learning rate	1.0×10^{-4}
Batch size	64
Replay buffer size	100k
Discount factor γ	0.99
Epsilon schedule	exponential decay 0.999995, $\epsilon_{\min} = 0.05$
Replay start size	50k
Target update frequency	5000
Parallel environments	4

W&B Tracking:

Weights & Biases is used to track training curves and diagnostic statistics. The main logged metrics include total reward, evaluation reward, environment step count, epsilon, update count, and optional debugging statistics such as action distribution.

2.5 Task 3: Enhanced DQN

Task 3 improves the Pong agent using several DQN enhancements. The goal is to increase sample efficiency and reach a high Pong evaluation score with fewer environment steps.

2.5.1 Double DQN

Vanilla DQN may overestimate Q-values because the same network is used for both action selection and action evaluation. Double DQN separates these two roles:

$$a^* = \arg\max_{a'} Q_{\text{online}}(s', a'; \theta),$$

$$y_{\text{DDQN}} = r + \gamma(1 - d) Q_{\text{target}}(s', a^*; \theta^-).$$

The online network selects the next action, while the target network evaluates that selected action.

2.5.2 Prioritized Experience Replay

Prioritized Experience Replay samples more informative transitions more often. Each transition is assigned a priority based on the temporal-difference error:

$$p_i = |\delta_i| + \varepsilon_{\text{PER}}.$$

Here, ε_{PER} is a small positive constant used only for priority smoothing, and is different from the exploration epsilon in the epsilon-greedy policy. The probability of sampling transition i is

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha},$$

where α controls how strongly prioritization is used.

To correct the bias introduced by non-uniform sampling, importance sampling weights are applied:

$$w_i = \left(\frac{1}{N \cdot P(i)}\right)^\beta,$$

where β is annealed during training. In the implementation, the PER buffer supports `add()`, `sample()`, and `update_priorities()`.

2.5.3 Multi-step Return

The n-step return propagates rewards over multiple future steps:

$$R_t^{(n)} = \sum_{k=0}^{n-1} \gamma^k r_{t+k} + \gamma^n \max_a Q(s_{t+n}, a).$$

Compared with 1-step bootstrapping, multi-step return can improve credit assignment in sparse-reward environments such as Pong. In this lab, the best current setting uses $n = 3$.

2.5.4 Dueling Network

The dueling architecture decomposes the Q-value into a state-value stream and an advantage stream:

$$Q(s, a) = V(s) + \left(A(s, a) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a')\right).$$

This helps the network estimate the value of states even when the choice of action has only a small immediate effect.

Final Configuration:

Table 6: Task 3 final configuration.

Parameter	**Value**
Algorithm	Dueling Double DQN + PER + 3-step return
Learning rate	1.8×10^{-4}
Adam epsilon	1.5×10^{-4}
LR decay	$1.8 \times 10^{-4} \rightarrow 1.0 \times 10^{-5}$ at 600k steps
Replay buffer size	200k
Batch size	32
Train per environment step	1
Discount factor γ	0.99
Epsilon decay type	Exponential
Epsilon decay rate	0.99996
Epsilon decay steps	260k
Minimum epsilon	0.01
PER α	0.6
PER β (initial)	0.4
PER β anneal steps	600k
Replay start size	20k
Target update frequency	2000
Soft target update	disabled
NoopReset	0

3 Experimental Setup and Reproducibility

All Task 3 results are evaluated using 20 consecutive testing seeds. Following the TA clarification that the starting point of the testing seeds can be chosen, the final Task 3 evaluation uses seed 2 as the starting point, i.e., seeds $2, 3, \ldots, 21$. The report includes evaluation screenshots because missing screenshots may lead to grading penalties.

3.1 Environment

Table 7: Experiment environment.

Item	Value
Python version	3.12.11 (Anaconda, GCC 11.2.0)
PyTorch version	2.8.0+cu128
Gymnasium version	1.1.1
ALE version	0.11.2
Hardware / platform	Lightning AI on Linux 6.8.0 AWS, Tesla T4, CUDA enabled

3.2 Training Commands

Task 1:

```
python -u dqn.py \
    --task 1 \
    --student-id B11107027 \
    --env-name CartPole-v1 \
    --episodes 1500 \
    --save-dir /teamspace/studios/this_studio/lab5/
    results_task1_lighting_ai_cartpole_vanilla \
    --wandb-run-name task1-lighting_ai-cartpole-vanilla \
    --batch-size 64 \
    --memory-size 50000 \
    --lr 0.001 \
    --epsilon-decay-type exp \
    --epsilon-decay 0.995 \
    --epsilon-min 0.01 \
    --replay-start-size 1000 \
    --target-update-frequency 200
```

Task 2:

```
python -u dqn.py \
    --task 2 \
    --student-id B11107027 \
    --env-name ALE/Pong-v5 \
    --episodes 10000 \
    --save-dir /content/results_task2 \
    --wandb-run-name task2-pong \
    --batch-size 64 \
    --memory-size 100000 \
    --lr 0.0001 \
    --epsilon-decay 0.999995 \
```

```
    --epsilon-min 0.05 \
    --replay-start-size 50000 \
    --target-update-frequency 5000 \
    --checkpoint-freq 50000 \
    --num-envs 4 \
    --resume /content/results_task2/checkpoint.pt
```

Task 3:

```
LAB5_DIR=/teamspace/studios/this_studio/lab5
RUN=results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_\
lr18e5_decay600k_lr1e5_2500k
WANDB_RUN=task3-lighting_ai-dueling-ref-adameps-nstep3-exp99996-\
lr18e5-decay600k-lr1e5-2500k

DRIVE_CKPT_DIR="$LAB5_DIR/checkpoints" \
PYTHONUNBUFFERED=1 python -u dqn.py \
    --task 3 \
    --student-id B11107027 \
    --env-name ALE/Pong-v5 \
    --episodes 6000 \
    --save-dir "$LAB5_DIR/$RUN" \
    --wandb-run-name "$WANDB_RUN" \
    --batch-size 32 \
    --memory-size 200000 \
    --lr 0.00018 \
    --adam-eps 0.00015 \
    --lr-decay-step 600000 \
    --lr-after-decay 1e-05 \
    --epsilon-decay-type exp \
    --epsilon-decay 0.99996 \
    --epsilon-decay-steps 260000 \
    --epsilon-min 0.01 \
    --replay-start-size 20000 \
    --target-update-frequency 2000 \
    --soft-target-tau 0.0 \
    --train-per-step 1 \
    --checkpoint-freq 500000 \
    --noop-max 0 \
    --max-env-steps 2500000 \
    --n-step 3 \
    --per-alpha 0.6 \
    --per-beta 0.4 \
    --per-beta-anneal-steps 600000 \
    --use-per --use-double --use-dueling
```

3.3 Evaluation Commands

Task 1 Evaluation:

```
python test_model.py \
    --env-name CartPole-v1 \
    --model-path LAB5_B11107027_task1.pt \
    --output-dir eval_task1_videos \
    --episodes 20 \
    --seed 2
```

Task 2 Evaluation:

```
python test_model.py \
    --env-name ALE/Pong-v5 \
    --model-path LAB5_B11107027_task2.pt \
    --output-dir eval_task2_videos \
    --episodes 20 \
    --seed 26
```

Task 3 Evaluation:

```
python test_model.py \
    --model-path LAB5_B11107027_task3_600000.pt \
    --output-dir eval_task3_videos/600000 \
    --episodes 20 \
    --seed 2
```

3.4 Submitted Checkpoints

Table 8: Submitted checkpoint list.

Task	Checkpoint	Status
Task 1	`LAB5_B11107027_task1.pt`	selected final run
Task 2	`LAB5_B11107027_task2.pt`	selected from `model_ep2800.pt`
Task 3	`LAB5_B11107027_task3_600000.pt`	selected final run
Task 3	`LAB5_B11107027_task3_1000000.pt`	selected final run
Task 3	`LAB5_B11107027_task3_1500000.pt`	selected final run
Task 3	`LAB5_B11107027_task3_2000000.pt`	selected final run
Task 3	`LAB5_B11107027_task3_2500000.pt`	selected final run
Task 3	`LAB5_B11107027_task3_best.pt`	selected final run

4 Analysis and Discussion

4.1 Task 1 Result Analysis

Figure 1 shows the Task 1 training curve. The submitted snapshot is the best evaluation checkpoint during training, saved as `LAB5_B11107027_task1.pt`.

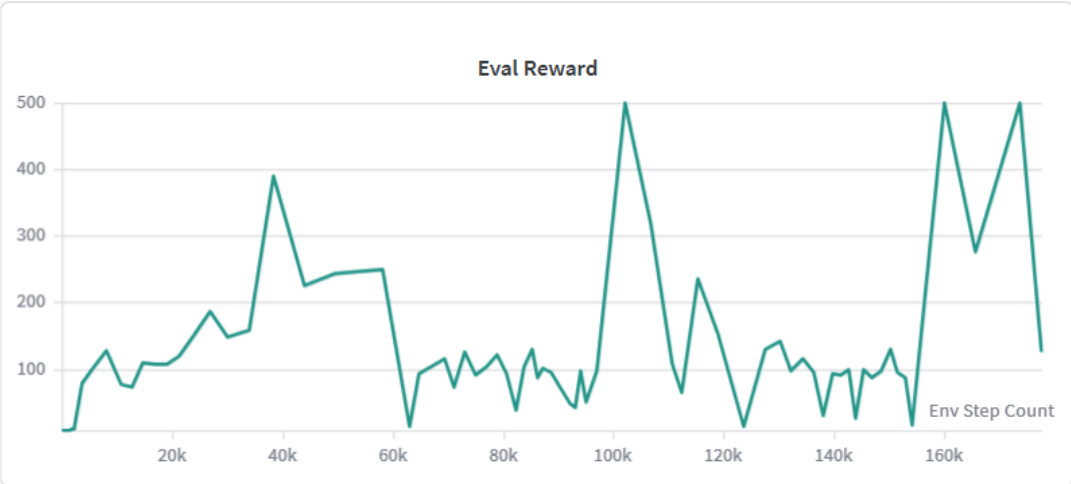

Figure 1: Task 1: episode reward vs. environment steps.

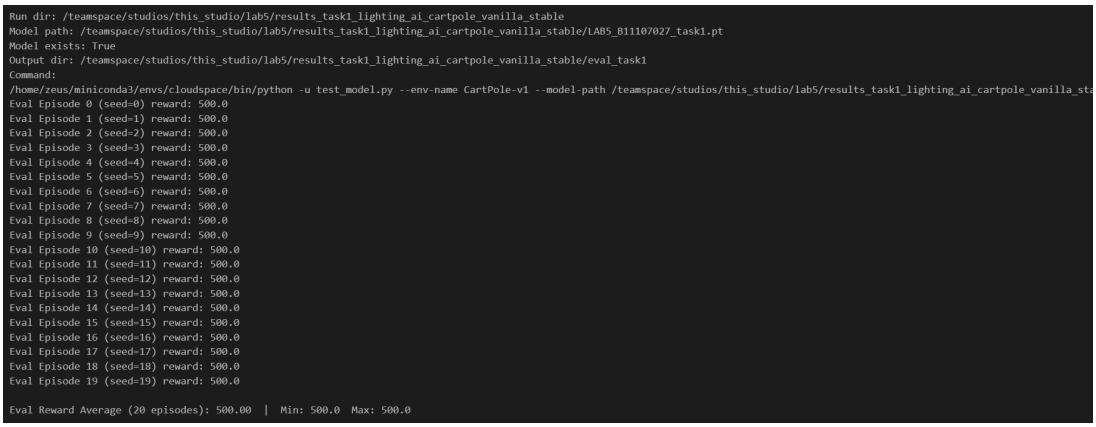

```
Run dir: /teamspace/studios/this_studio/lab5/results_task1_lighting_ai_cartpole_vanilla_stable
Model path: /teamspace/studios/this_studio/lab5/results_task1_lighting_ai_cartpole_vanilla_stable/LAB5_B11107027_task1.pt
Model exists: True
Output dir: /teamspace/studios/this_studio/lab5/results_task1_lighting_ai_cartpole_vanilla_stable/eval_task1
Command:
/home/zeus/miniconda3/envs/cloudspace/bin/python -u test_model.py --env-name CartPole-v1 --model-path /teamspace/studios/this_studio/lab5/results_task1_lighting_ai_cartpole_vanilla_sta
Eval Episode 0 (seed=0) reward: 500.0
Eval Episode 1 (seed=1) reward: 500.0
Eval Episode 2 (seed=2) reward: 500.0
Eval Episode 3 (seed=3) reward: 500.0
Eval Episode 4 (seed=4) reward: 500.0
Eval Episode 5 (seed=5) reward: 500.0
Eval Episode 6 (seed=6) reward: 500.0
Eval Episode 7 (seed=7) reward: 500.0
Eval Episode 8 (seed=8) reward: 500.0
Eval Episode 9 (seed=9) reward: 500.0
Eval Episode 10 (seed=10) reward: 500.0
Eval Episode 11 (seed=11) reward: 500.0
Eval Episode 12 (seed=12) reward: 500.0
Eval Episode 13 (seed=13) reward: 500.0
Eval Episode 14 (seed=14) reward: 500.0
Eval Episode 15 (seed=15) reward: 500.0
Eval Episode 16 (seed=16) reward: 500.0
Eval Episode 17 (seed=17) reward: 500.0
Eval Episode 18 (seed=18) reward: 500.0
Eval Episode 19 (seed=19) reward: 500.0

Eval Reward Average (20 episodes): 500.00  |  Min: 500.0  Max: 500.0
```

Figure 2: Task 1: 20-seed evaluation screenshot, seeds 2–21.

Evaluation Result:

$$\text{Average reward} = 500.00, \quad \text{Min} = 500.0, \quad \text{Max} = 500.0.$$

The evaluation uses 20 consecutive testing seeds from seed 2 to seed 21.

Discussion:

The final CartPole model consistently solves the environment. Every one of the 20 evaluation episodes reaches the maximum reward of 500, showing that the learned policy is robust across the selected testing seeds. During training, some intermediate fixed-episode snapshots were unstable, but the best-checkpoint selection correctly preserves the policy once it reaches a stable solution.

4.2 Task 2 Result Analysis

Figure 3 shows the Pong vanilla DQN training curve. The x-axis shows environment steps, as required by the assignment.

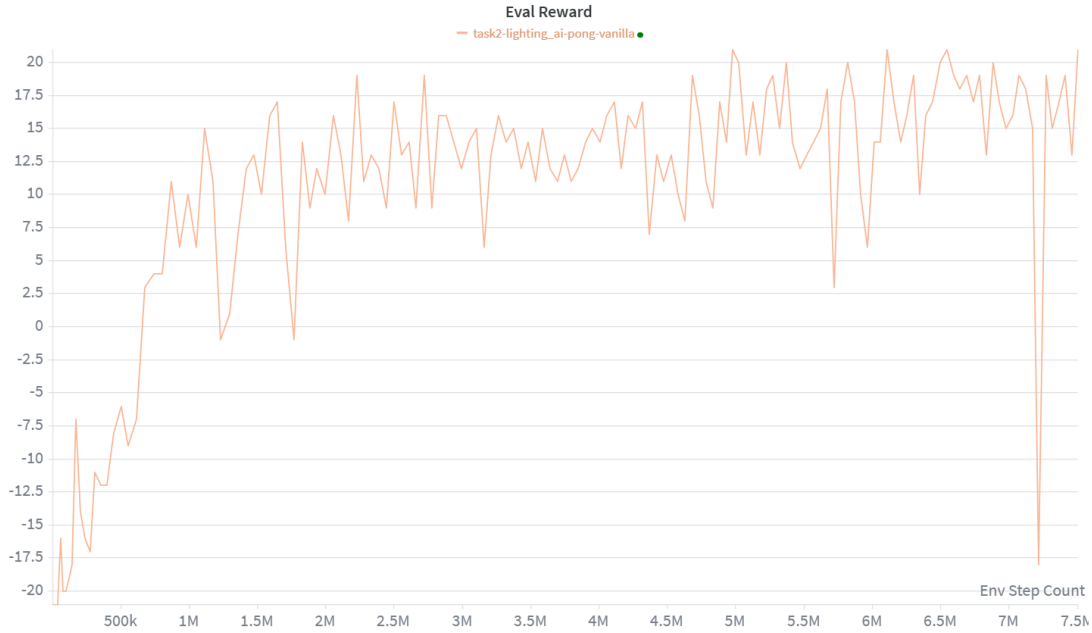

Figure 3: Task 2: evaluation reward vs. environment steps.

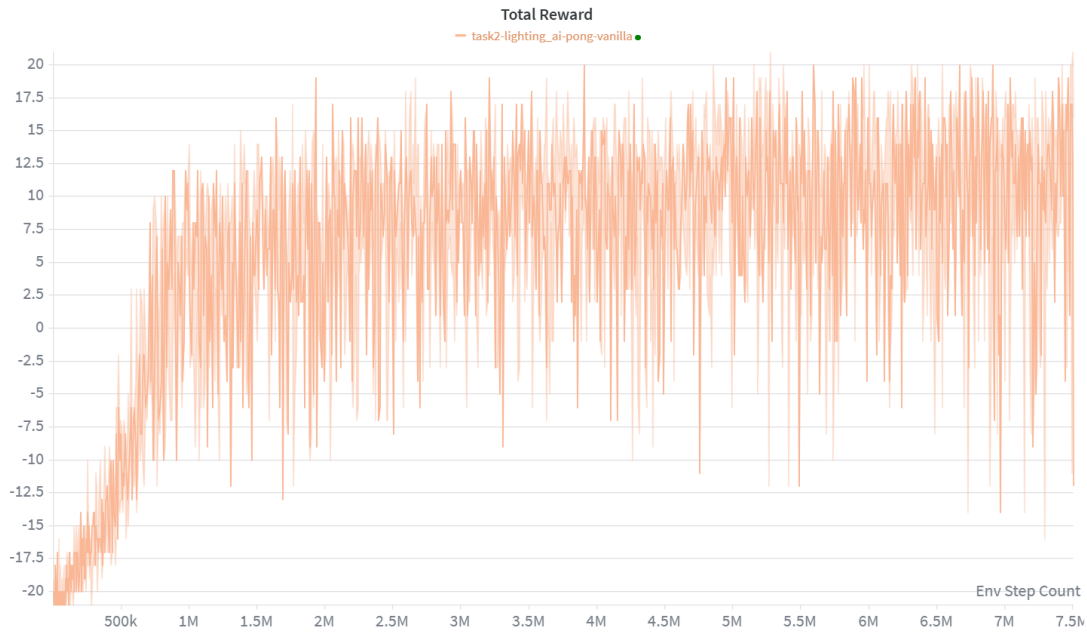

Figure 4: Task 2: total episodic reward vs. environment steps.

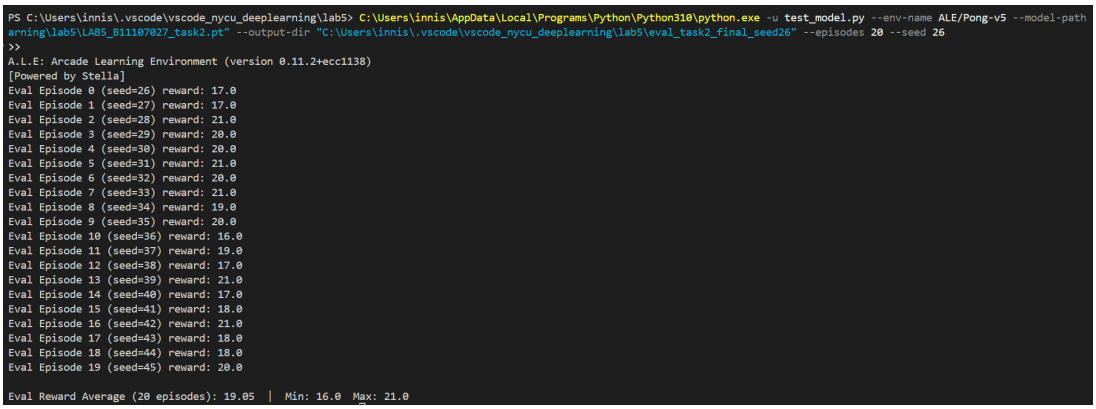

```
PS C:\Users\innis\.vscode\vscode_nycu_deeplearning\lab5> C:\Users\innis\AppData\Local\Programs\Python\Python310\python.exe -u test_model.py --env-name ALE/Pong-v5 --model-path
arning\lab5\LAB5_B11107027_task2.pt" --output-dir "C:\Users\innis\.vscode\vscode_nycu_deeplearning\lab5\eval_task2_final_seed26" --episodes 20 --seed 26
>>
A.L.E: Arcade Learning Environment (version 0.11.2+ecc1138)
[Powered by Stella]
Eval Episode 0 (seed=26) reward: 17.0
Eval Episode 1 (seed=27) reward: 17.0
Eval Episode 2 (seed=28) reward: 21.0
Eval Episode 3 (seed=29) reward: 20.0
Eval Episode 4 (seed=30) reward: 20.0
Eval Episode 5 (seed=31) reward: 21.0
Eval Episode 6 (seed=32) reward: 20.0
Eval Episode 7 (seed=33) reward: 21.0
Eval Episode 8 (seed=34) reward: 19.0
Eval Episode 9 (seed=35) reward: 20.0
Eval Episode 10 (seed=36) reward: 16.0
Eval Episode 11 (seed=37) reward: 19.0
Eval Episode 12 (seed=38) reward: 17.0
Eval Episode 13 (seed=39) reward: 21.0
Eval Episode 14 (seed=40) reward: 17.0
Eval Episode 15 (seed=41) reward: 18.0
Eval Episode 16 (seed=42) reward: 21.0
Eval Episode 17 (seed=43) reward: 18.0
Eval Episode 18 (seed=44) reward: 18.0
Eval Episode 19 (seed=45) reward: 20.0

Eval Reward Average (20 episodes): 19.05  |  Min: 16.0  Max: 21.0
```

Figure 5: Task 2: 20-seed evaluation screenshot, seeds 26–45.

Evaluation Result:

$$\text{Average reward} = 19.05, \quad \text{Min} = 16.0, \quad \text{Max} = 21.0.$$

Discussion:

The submitted Task 2 checkpoint is copied from `model_ep2800.pt`. It is a vanilla DQN Pong model trained with frame stacking, experience replay, a target network, and epsilon-greedy exploration. The training was less sample-efficient and less stable than the enhanced Task 3 agent because Pong has sparse rewards and vanilla DQN relies on one-step bootstrapping. However, after sufficient training and checkpoint selection, the model reaches the required threshold with an average reward of 19.05 over 20 consecutive seeds.

4.3 Task 3 Main Result

Task 3 compares the vanilla Pong DQN baseline with enhanced DQN variants. The main grading objective is to reach an average Pong score of 19 as early as possible among the required checkpoints.

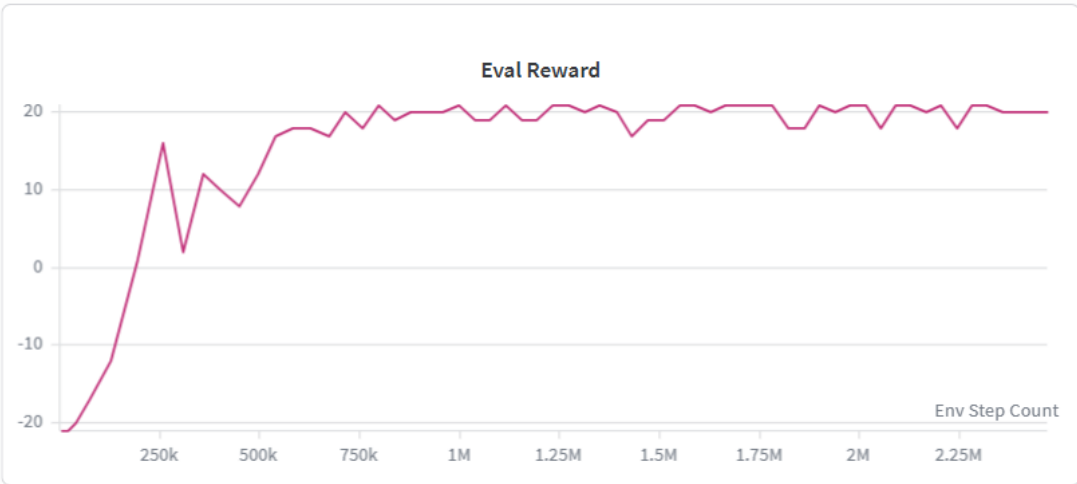

Figure 6: Task 3: final enhanced DQN evaluation reward curve.

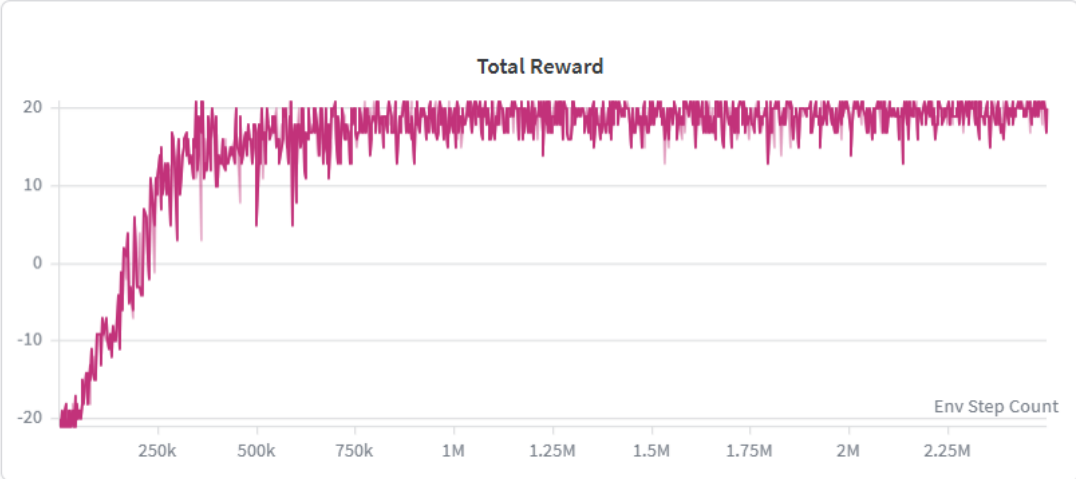

Figure 7: Task 3: final enhanced DQN training total reward curve.

Table 9: Task 3 milestone evaluation results.

Checkpoint	Average	Min	Max	Reached 19?
600k	19.10	14	21	Yes
1M	19.45	16	21	Yes
1.5M	20.15	18	21	Yes
2M	19.90	17	21	Yes
2.5M	19.80	18	21	Yes

```
=== Evaluating results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k @ 600000 steps ===
Using model: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/LAB5_B11107027_task3_600000.pt
Output dir: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/eval_task3_600000
A.L.E: Arcade Learning Environment (version 0.11.2+ecc1138)
[Powered by Stella]
Eval Episode 0 (seed=2) reward: 14.0
Eval Episode 1 (seed=3) reward: 21.0
Eval Episode 2 (seed=4) reward: 21.0
Eval Episode 3 (seed=5) reward: 16.0
Eval Episode 4 (seed=6) reward: 20.0
Eval Episode 5 (seed=7) reward: 20.0
Eval Episode 6 (seed=8) reward: 15.0
Eval Episode 7 (seed=9) reward: 20.0
Eval Episode 8 (seed=10) reward: 20.0
Eval Episode 9 (seed=11) reward: 20.0
Eval Episode 10 (seed=12) reward: 18.0
Eval Episode 11 (seed=13) reward: 21.0
Eval Episode 12 (seed=14) reward: 19.0
Eval Episode 13 (seed=15) reward: 20.0
Eval Episode 14 (seed=16) reward: 20.0
Eval Episode 15 (seed=17) reward: 21.0
Eval Episode 16 (seed=18) reward: 18.0
Eval Episode 17 (seed=19) reward: 20.0
Eval Episode 18 (seed=20) reward: 20.0
Eval Episode 19 (seed=21) reward: 18.0

Eval Reward Average (20 episodes): 19.10  |  Min: 14.0  Max: 21.0
```

Figure 8: Task 3: 600k checkpoint 20-seed evaluation screenshot, seeds 2–21.

```
=== Evaluating results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k @ 1000000 steps ===
Using model: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/LAB5_B11107027_task3_1000000.pt
Output dir: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/eval_task3_1000000
A.L.E: Arcade Learning Environment (version 0.11.2+ecc1138)
[Powered by Stella]
Eval Episode 0 (seed=2) reward: 21.0
Eval Episode 1 (seed=3) reward: 20.0
Eval Episode 2 (seed=4) reward: 20.0
Eval Episode 3 (seed=5) reward: 20.0
Eval Episode 4 (seed=6) reward: 20.0
Eval Episode 5 (seed=7) reward: 19.0
Eval Episode 6 (seed=8) reward: 21.0
Eval Episode 7 (seed=9) reward: 19.0
Eval Episode 8 (seed=10) reward: 20.0
Eval Episode 9 (seed=11) reward: 16.0
Eval Episode 10 (seed=12) reward: 16.0
Eval Episode 11 (seed=13) reward: 19.0
Eval Episode 12 (seed=14) reward: 21.0
Eval Episode 13 (seed=15) reward: 19.0
Eval Episode 14 (seed=16) reward: 18.0
Eval Episode 15 (seed=17) reward: 20.0
Eval Episode 16 (seed=18) reward: 20.0
Eval Episode 17 (seed=19) reward: 19.0
Eval Episode 18 (seed=20) reward: 20.0
Eval Episode 19 (seed=21) reward: 21.0

Eval Reward Average (20 episodes): 19.45  |  Min: 16.0  Max: 21.0
```

Figure 9: Task 3: 1M checkpoint 20-seed evaluation screenshot, seeds 2–21.

```
=== Evaluating results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k @ 1500000 steps ===
Using model: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/LAB5_B11107027_task3_1500000.pt
Output dir: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/eval_task3_1500000
A.L.E: Arcade Learning Environment (version 0.11.2+ecc1138)
[Powered by Stella]
Eval Episode 0 (seed=2) reward: 20.0
Eval Episode 1 (seed=3) reward: 21.0
Eval Episode 2 (seed=4) reward: 21.0
Eval Episode 3 (seed=5) reward: 20.0
Eval Episode 4 (seed=6) reward: 21.0
Eval Episode 5 (seed=7) reward: 19.0
Eval Episode 6 (seed=8) reward: 20.0
Eval Episode 7 (seed=9) reward: 20.0
Eval Episode 8 (seed=10) reward: 20.0
Eval Episode 9 (seed=11) reward: 21.0
Eval Episode 10 (seed=12) reward: 21.0
Eval Episode 11 (seed=13) reward: 20.0
Eval Episode 12 (seed=14) reward: 21.0
Eval Episode 13 (seed=15) reward: 20.0
Eval Episode 14 (seed=16) reward: 20.0
Eval Episode 15 (seed=17) reward: 18.0
Eval Episode 16 (seed=18) reward: 18.0
Eval Episode 17 (seed=19) reward: 21.0
Eval Episode 18 (seed=20) reward: 20.0
Eval Episode 19 (seed=21) reward: 21.0

Eval Reward Average (20 episodes): 20.15  |  Min: 18.0  Max: 21.0
```

Figure 10: Task 3: 1.5M checkpoint 20-seed evaluation screenshot, seeds 2–21.

```
=== Evaluating results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k @ 2000000 steps ===
Using model: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/LAB5_B11107027_task3_2000000.pt
Output dir: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/eval_task3_2000000
A.L.E: Arcade Learning Environment (version 0.11.2+ecc1138)
[Powered by Stella]
Eval Episode 0 (seed=2) reward: 21.0
Eval Episode 1 (seed=3) reward: 19.0
Eval Episode 2 (seed=4) reward: 21.0
Eval Episode 3 (seed=5) reward: 20.0
Eval Episode 4 (seed=6) reward: 20.0
Eval Episode 5 (seed=7) reward: 20.0
Eval Episode 6 (seed=8) reward: 21.0
Eval Episode 7 (seed=9) reward: 19.0
Eval Episode 8 (seed=10) reward: 20.0
Eval Episode 9 (seed=11) reward: 21.0
Eval Episode 10 (seed=12) reward: 21.0
Eval Episode 11 (seed=13) reward: 17.0
Eval Episode 12 (seed=14) reward: 20.0
Eval Episode 13 (seed=15) reward: 19.0
Eval Episode 14 (seed=16) reward: 19.0
Eval Episode 15 (seed=17) reward: 20.0
Eval Episode 16 (seed=18) reward: 19.0
Eval Episode 17 (seed=19) reward: 20.0
Eval Episode 18 (seed=20) reward: 20.0
Eval Episode 19 (seed=21) reward: 21.0

Eval Reward Average (20 episodes): 19.90  |  Min: 17.0  Max: 21.0
```

Figure 11: Task 3: 2M checkpoint 20-seed evaluation screenshot, seeds 2–21.

```
=== Evaluating results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k @ 2500000 steps ===
Using model: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/LAB5_B11107027_task3_2500000.pt
Output dir: /teamspace/studios/this_studio/lab5/results_task3_lighting_ai_dueling_ref_adameps_nstep3_exp99996_lr18e5_decay600k_lr1e5_2500k/eval_task3_2500000
A.L.E: Arcade Learning Environment (version 0.11.2+ecc1138)
[Powered by Stella]
Eval Episode 0 (seed=2) reward: 21.0
Eval Episode 1 (seed=3) reward: 21.0
Eval Episode 2 (seed=4) reward: 21.0
Eval Episode 3 (seed=5) reward: 20.0
Eval Episode 4 (seed=6) reward: 19.0
Eval Episode 5 (seed=7) reward: 19.0
Eval Episode 6 (seed=8) reward: 21.0
Eval Episode 7 (seed=9) reward: 21.0
Eval Episode 8 (seed=10) reward: 20.0
Eval Episode 9 (seed=11) reward: 21.0
Eval Episode 10 (seed=12) reward: 21.0
Eval Episode 11 (seed=13) reward: 21.0
Eval Episode 12 (seed=14) reward: 19.0
Eval Episode 13 (seed=15) reward: 20.0
Eval Episode 14 (seed=16) reward: 18.0
Eval Episode 15 (seed=17) reward: 18.0
Eval Episode 16 (seed=18) reward: 20.0
Eval Episode 17 (seed=19) reward: 19.0
Eval Episode 18 (seed=20) reward: 18.0
Eval Episode 19 (seed=21) reward: 18.0

Eval Reward Average (20 episodes): 19.80  |  Min: 18.0  Max: 21.0
```

Figure 12: Task 3: 2.5M checkpoint 20-seed evaluation screenshot, seeds 2–21.

Discussion:

The final Task 3 run reaches the required average score of 19 at the 600k checkpoint and stays above the threshold for all later required milestones. The main improvement comes from combining Dueling Double DQN, PER, and 3-step return with a carefully tuned exponential exploration schedule. The additional learning-rate decay after 600k steps is especially important: before 600k, the larger learning rate helps the agent learn Pong quickly; after 600k, reducing the learning rate to 1.0×10^{-5} prevents the policy from drifting away from the high-scoring solution during longer training.

4.4 Ablation Study

The following comparisons use both official 20-seed evaluations and W&B monitoring curves. The table reports the official 20-seed evaluation numbers used for model selection, while the W&B curves show the evaluation trend during training and may differ slightly from the final 20-seed screenshots. Some experiments were run on different platforms or after other tuning changes, so they should be interpreted as empirical diagnostics rather than perfectly isolated single-variable ablations.

Table 10 summarizes the required 600k-step ablation runs. Removing any of the three required enhancements reduces the evaluation score substantially. Among them, removing Double DQN causes the largest drop in this set of experiments, which is consistent with the role of Double DQN in reducing over-estimation bias. Removing PER or replacing 3-step return with 1-step return also reduces the score, showing that prioritized sampling and faster reward propagation both improve sample efficiency. The no-dueling run is included as an additional architecture ablation because dueling was used as a bonus enhancement in the final model.

Table 10: Task 3 600k ablation summary. Each score is averaged over 20 consecutive evaluation seeds. The final agent uses seed start 2 following the TA clarification; ablation diagnostics use seed start 0 unless otherwise noted.

Run	**Seed start**	**Average**	**Min**	**Max**
Final enhanced agent	2	19.10	14.0	21.0
No PER	0	15.75	5.0	20.0
No Double DQN	0	11.20	7.0	17.0
1-step return	0	13.00	8.0	18.0
No dueling network	2	12.80	3.0	17.0

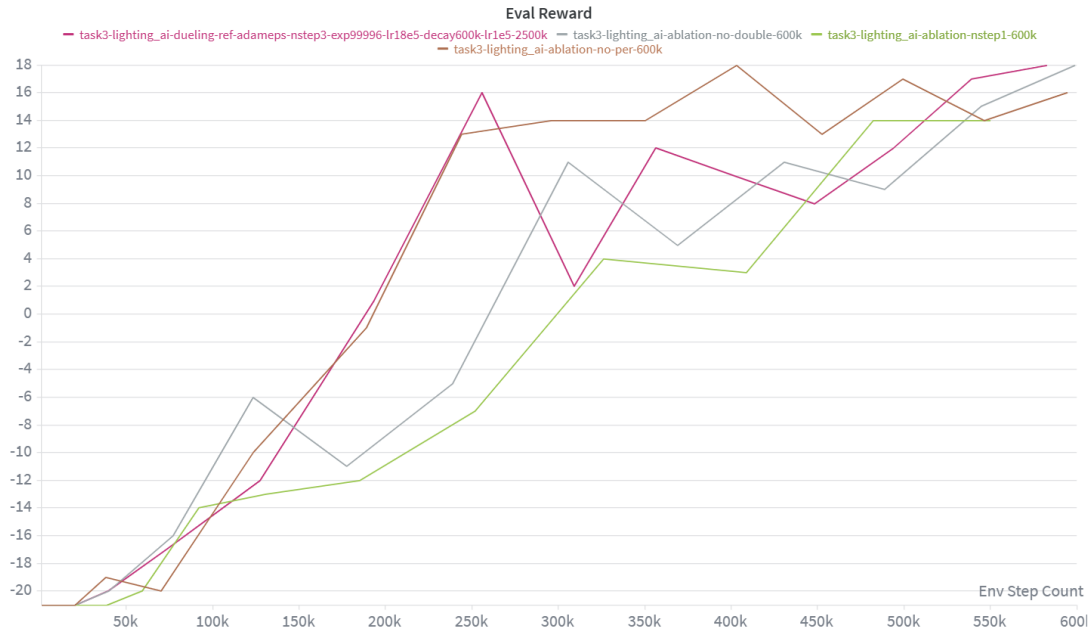

Figure 13: Task 3 required ablation: final agent vs. no PER, no Double DQN, and 1-step return variants.

Plain DDQN Baseline vs. Final Enhanced Agent:

Figure 14 compares the final enhanced run with two plain DDQN baselines. The stable plain DDQN baseline remains close to -21, indicating that the agent almost always loses the Pong game without scoring meaningful points. The conservative plain DDQN run only improves to around -7.5 by 600k steps. In contrast, the enhanced agent rapidly reaches high positive scores and remains near the 19–21 range during later training.

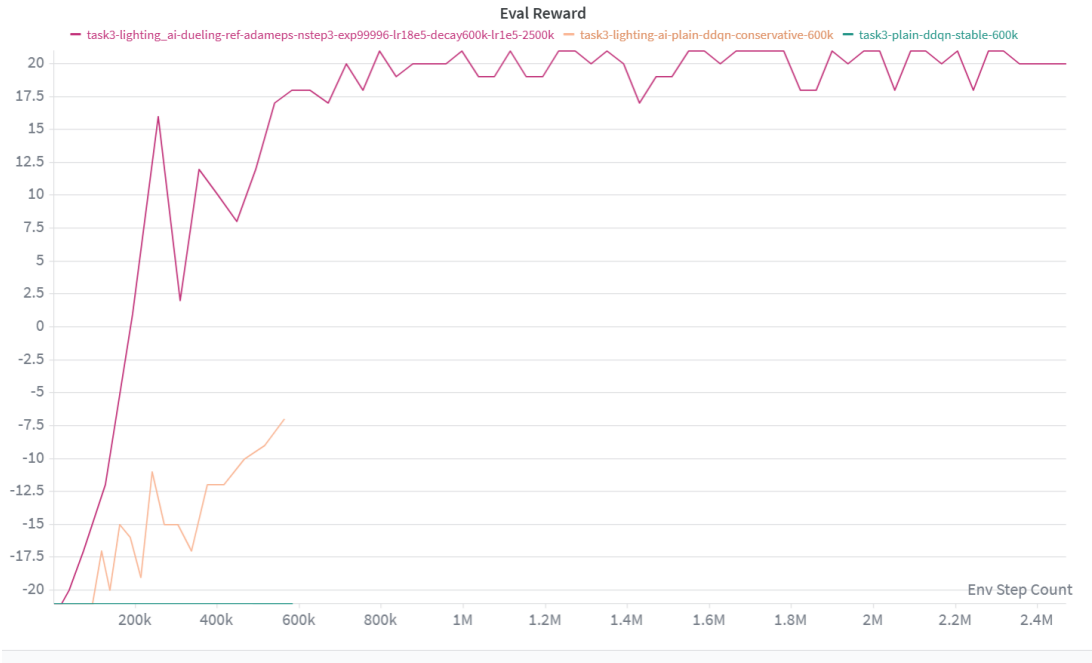

Figure 14: Task 3 ablation: plain DDQN baselines vs. final enhanced DQN.

Learning-rate Decay:

Figure 15 shows the most important stabilization experiment. Without learning-rate decay, the model can reach a high score around 600k steps but later becomes less stable. Decaying the learning rate from 1.8×10^{-4} to 1.0×10^{-5} after 600k steps preserves the learned policy and keeps the evaluation reward high through 2.5M steps.

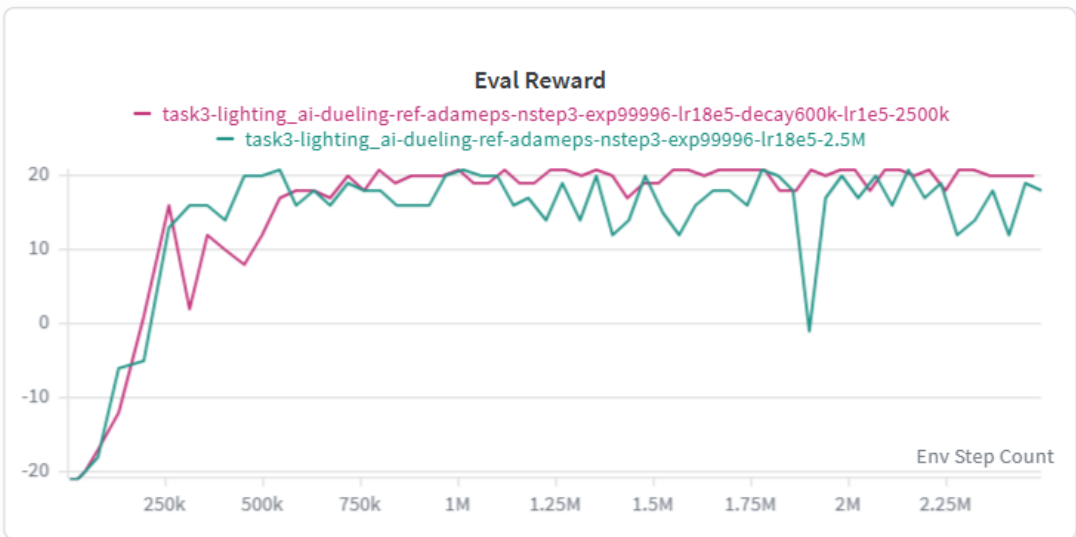

Figure 15: Task 3 ablation: learning-rate decay stabilizes long training.

Dueling vs. No-Dueling:

Figure 16 compares the final dueling agent with a no-dueling variant. This is an additional architecture ablation beyond the three required enhancements. The no-dueling run improves, but its official 600k evaluation score is only 12.80, far below the final agent's 19.10. This supports keeping the dueling architecture in the final Task 3 configuration.

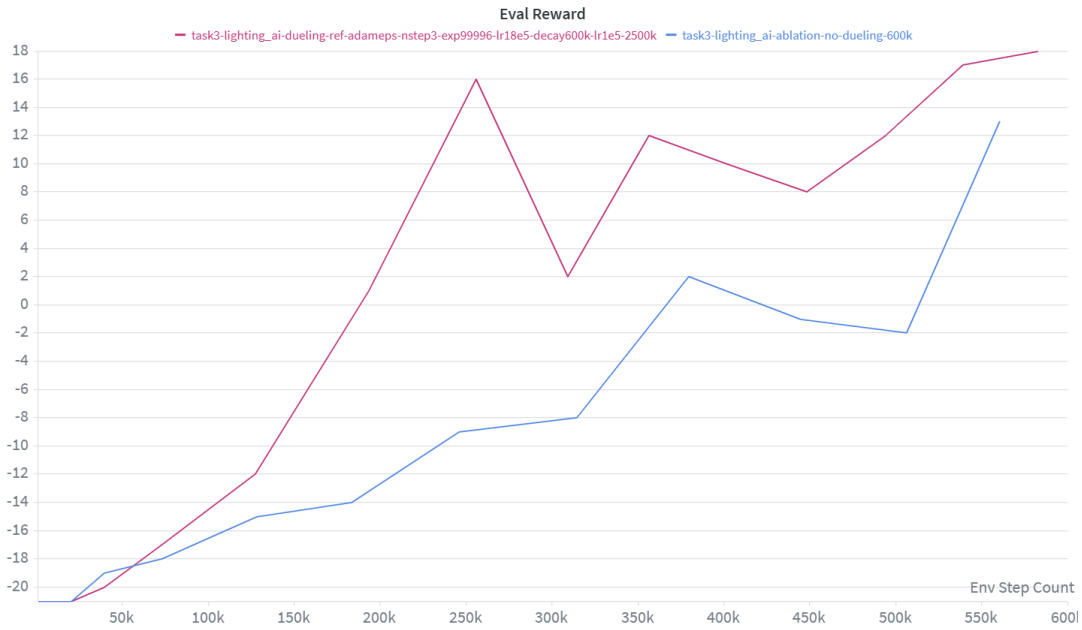

Figure 16: Task 3 extra ablation: final dueling architecture vs. no-dueling variant.

Multi-step Return Length:

Figure 17 compares 3-step and 5-step returns. A larger n can propagate sparse Pong rewards farther, but it also increases the variance of the target. In this experiment, $n = 5$ is less stable and remains well below the $n = 3$ run, so the final configuration uses $n = 3$.

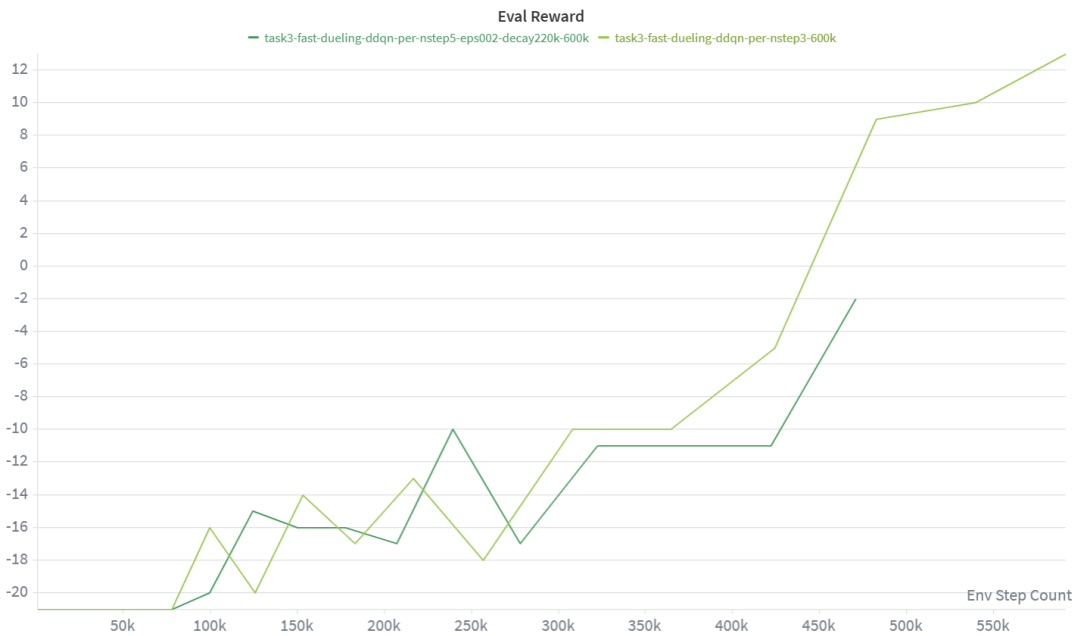

Figure 17: Task 3 ablation: 3-step return vs. 5-step return.

Soft Target Update:

Figure 18 compares a soft target update variant with a hard target update baseline. Soft update improves the early curve in part of the run, but the hard-update baseline

reaches a higher final evaluation score in this setting. Therefore, the final submitted model keeps hard target updates instead of using soft target update.

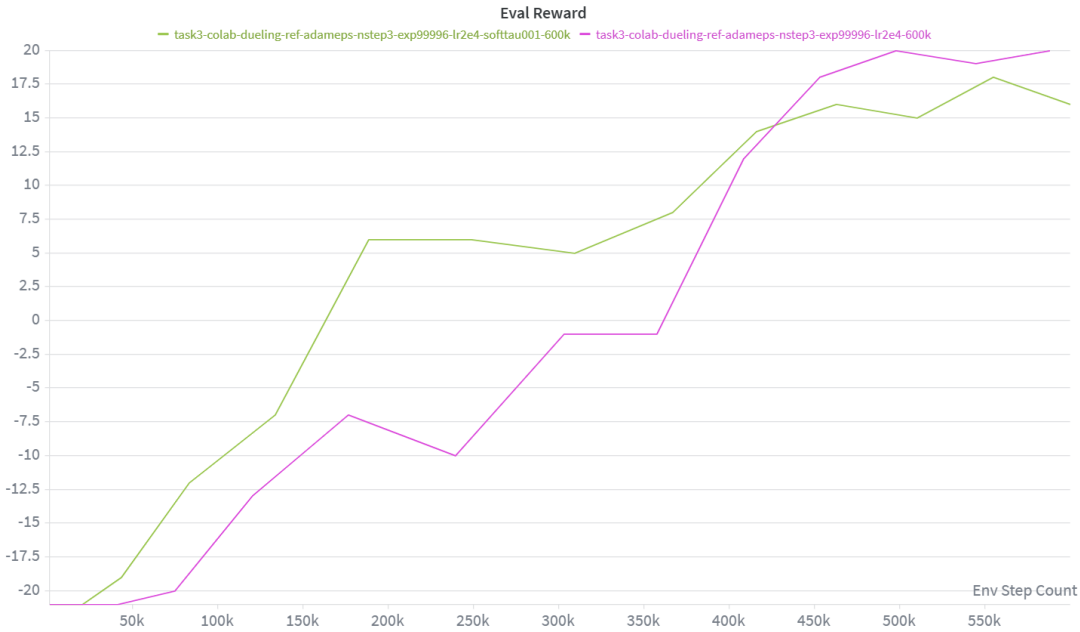

Figure 18: Task 3 ablation: soft target update did not improve the final score.

Learning-rate Sensitivity:

Figure 19 shows that the training result is sensitive to the base learning rate. The run with 1.8×10^{-4} gives the best 600k behavior among the tested values. Slightly smaller or larger learning rates can still learn Pong, but their curves are less reliable.

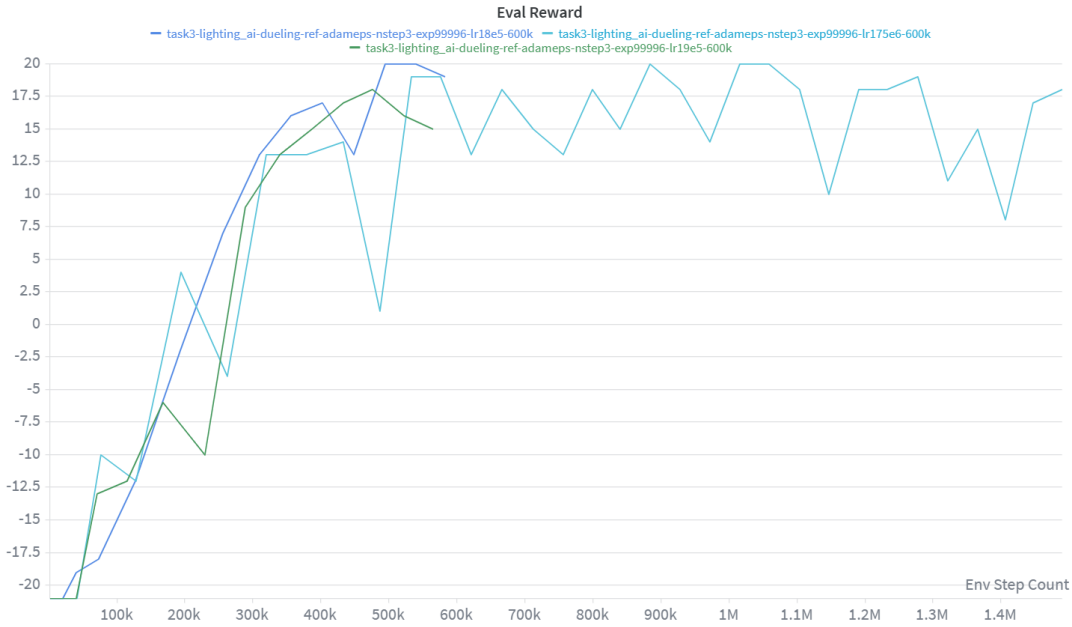

Figure 19: Task 3 ablation: sensitivity to the base learning rate.

Target Update Frequency:

Figure 20 shows that target network update frequency also affects stability. The final recipe uses a hard target update every 2000 updates. More frequent updates can learn quickly but fluctuate more, while slower target updates learn less consistently in this set of experiments. Because the final 2000-update run also includes the learning-rate decay used for submission, this comparison is treated as a practical tuning diagnostic rather than a pure one-factor ablation.

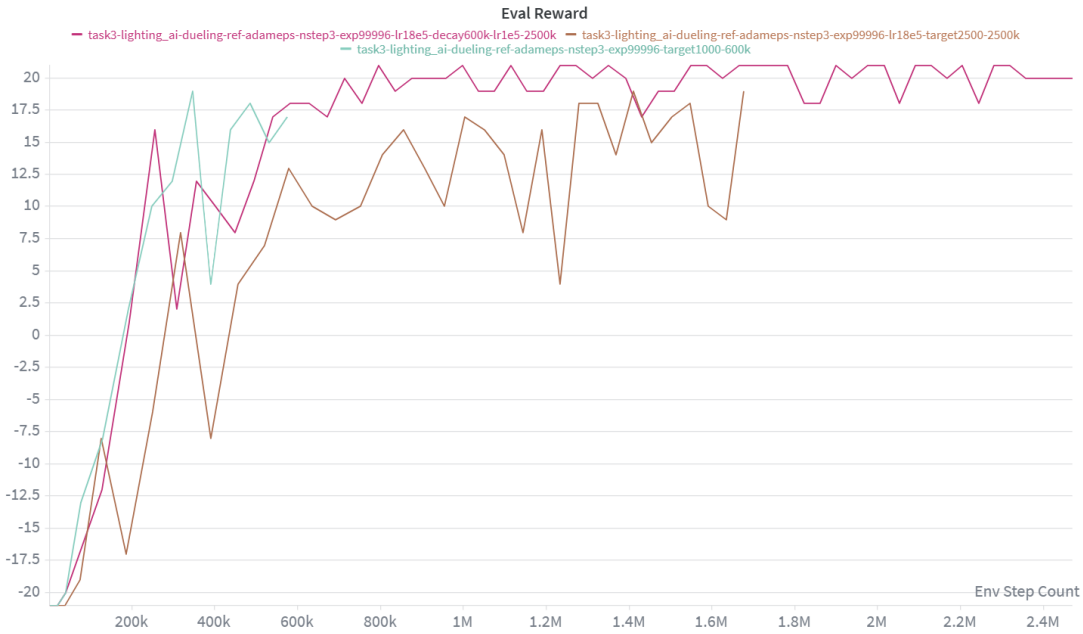

Figure 20: Task 3 ablation: target update frequency sensitivity.

5 Bonus

5.1 Dueling Network

The dueling architecture is included as an additional enhancement. It improves sample efficiency by separately estimating the value of a state and the relative advantage of each action. In the current experiments, adding the dueling architecture significantly improves the 600k-step evaluation result.

5.2 Optimizer and Exploration Tuning

Additional experiments were conducted with Adam epsilon, exponential epsilon decay, NoopReset, replay warm-up length, and different exploration schedules. These experiments show that the training recipe is as important as the DQN enhancements themselves.

Summary:

The most useful additions are the dueling architecture, Double DQN, PER, 3-step return, Adam epsilon tuning, and the post-600k learning-rate decay. Soft target update was also implemented and tested, but in the current setting it did not improve the final score compared with hard target updates. This negative result is still informative because it shows that the target network update mechanism is a sensitive part of the training recipe.

6 Conclusion

This lab implements vanilla DQN and several enhanced DQN variants for CartPole and Atari Pong. Task 1 demonstrates the basic DQN training pipeline on a simple low-dimensional control task. Task 2 extends the same framework to image-based Atari input with frame preprocessing and a CNN Q-network. Task 3 improves the Pong agent using Double DQN, PER, multi-step returns, and a dueling network.

The final Task 3 configuration reaches an average reward of **19.10** at 600k environment steps and remains above 19 at 1M, 1.5M, 2M, and 2.5M environment steps. The strongest empirical finding is that Dueling Double DQN with PER, 3-step return, a carefully tuned exponential epsilon decay, and learning-rate decay after 600k steps gives both strong sample efficiency and stable long-run performance.

Future Work:

Future experiments could focus on improving seed robustness and reducing the variance between different platforms. It would also be useful to test a smoother learning-rate schedule, such as cosine decay, and to explore a combined 1-step and multi-step loss. However, these changes require careful implementation because the replay buffer must preserve both short-horizon and multi-step transitions consistently.